

UNIVERSITY OF OTAGO

DEPARTMENT OF COMPUTER SCIENCE

COSC480 PROJECT REPORT

Quantifying Conceptual Density in Text

Author:

Anthony DICKSON
(3348967)

Supervisor(s):

Anthony ROBINS
Alistair KNOTT

July 22, 2019



Abstract

Conceptual density is an idea related to the degree to which concepts in a domain are integrated, or interdependent. There is a hypothesis that text documents with high conceptual density are harder to process. In this project, a definition of conceptual density is given and a graph-based method for quantifying conceptual density in well-structured expository text documents, such as educational textbooks, are explored.

1 Introduction

[Robins \(2010\)](#) suggested that a domain of tightly integrated concepts magnifies the effect of early success or failure of learning, leading to polarised learning outcomes, an effect that he named learning edge momentum. A tightly integrated concept is said to be a concept that is difficult to describe or explain independently of other concepts. This project is focused on the idea of conceptual density, an idea related to the degree to which concepts in a given domain are integrated.

The idea of conceptual density affecting the difficulty of learning is not unique to the aforementioned work. Cognitive load theory provides a framework for reasoning about aspects of learning difficulty and can be used to guide instructional design ([Sweller, 1994](#)). Conceptual density appears in the literature on cognitive load theory, in a perhaps slightly different form, as “element interactivity”. Element interactivity, similar to conceptual density, is an idea related to the number of interacting elements in a task or piece of material that is to be learned. Under the element interactivity effect, high element interactivity is suggested to be the main contributor to the difficulty of a learning task.

To the best of my knowledge, there exists no objective measure in the current literature of the element interactivity or the notion of conceptual density. Element interactivity appears to be estimated manually and it depends greatly on the knowledge of the learner ([Chandler and Sweller, 1996](#)), and as such remains subjective. In ([Robins, 2010](#)) the notion of concepts being tightly integrated stems from the observation that educators typically fail to reach a consensus on how to structure introductory programming courses (this is supposedly not the case in many other disciplines). While no measure of conceptual density was clearly defined in that work, some possible ways to do so were proposed.

In this project, I follow up on the notion of conceptual density and set out to elaborate on its definition and develop an objective measure of conceptual density in the context of well-structured expository text documents (e.g. textbooks). For measuring conceptual density I explore one of the methods mentioned in (Robins, 2010) - analysing a mind-map like graph structure. Whereas in the original work, this approach was suggested as a characterisation of introductory programming course material that is to be generated by experts, in this project I aim to build up and analyse mind-map like graph structures through the use of natural language processing methods and graph analysis techniques, removing the need for an expert.

The rest of this report is structured as follows: Section 2 covers the theoretical background of conceptual density; Section 3 defines what is considered as a concept in text and other text features that relate to conceptual density; Section 4 describes the parsing algorithm used for building a mind-map like graph structure and how a numerical score of conceptual density can be derived from a graph structure; Section 5 discusses related work and what distinguishes the work in this project from the related work; and Section 6 discusses the work-to-date and future work.

2 Background

2.1 Learning Edge Momentum

Robins (2010) proposed learning edge momentum as an explanation for the unusually high rate of both fail and high grades in introductory programming courses. The hypothesis builds upon the idea that we learn at the edges of what we know, or in other words, it is easy to learn things that build upon pre-existing knowledge. The core of the hypothesis is that learning outcomes are self-reinforcing: successful learning of a concept tends to make learning closely related concepts easier, and unsuccessful learning tends to make learning closely related concepts more difficult. It is also said that concepts in programming languages tend to be tightly integrated and that these tightly integrated concepts are difficult to describe or understand independently of each other. This means that learning a new concept is greatly dependent on understanding the prerequisite concepts. And because of this, the effects of learning edge momentum in introductory programming courses are particularly pronounced, leading to the polarised distribution of grades.

Conceptual density is mentioned in this paper as a possible metric for providing more concrete evidence to educators on how to structure introductory programming courses. Two methods are suggested for measuring the degree to which concepts in a domain are integrated. One method is to look at the proximity of key concept terms and see if there are areas where many concepts are mentioned close together. The other method is to generate a mind-map like graph structure. In this approach, a domain of concepts is represented as a graph structure, where the concepts present in the material are represented with a node in the graph and edges between nodes represent a connection between concepts. This follows the main ideas behind semantic networks and knowledge graphs (Sowa, 1987; Zhang, 2002; Koncel-Kedziorski et al., 2019), which are graph structures used for knowledge representation. This idea of a mind map of concepts forms the basis for the work in this project.

2.2 Cognitive Load Theory and Schema

CLT (cognitive load theory) provides a framework for reasoning about learning difficulty and instructional design (Sweller, 1994). In CLT the main task of the brain is characterised as information processing, and each person has a certain amount of working memory to process information. Performing tasks requires an amount of mental exertion and concentration, called cognitive load. When the amount of cognitive load exceeds the capacity of an individual’s working memory, performing the given task becomes much more difficult.

If we are to explain this by analogy, the idea of working memory is similar in some ways to computer RAM and running out of working memory is akin to what happens when a computer runs out of RAM. When a computer runs out of RAM, it starts thrashing and all of the programs will become very slow and take many times longer to run compared to when there is enough RAM. By reducing the number of programs that are running we can reduce the memory usage and prevent the computer from thrashing. Similarly, under CLT we can reduce cognitive load by reducing the number of things that must be considered simultaneously to perform a given task in order to make it easier to perform that task.

The main contributor to the difficulty of a task is the number of things that must be considered simultaneously and the degree to which they interact. In CLT this idea is captured by **element interactivity** and the

```

1 public class HelloWorld {
2     public static void main(String[] args) {
3         // Prints "Hello, World" to the terminal window.
4         System.out.println("Hello, World");
5     }
6 }

```

Figure 1: Example Java program that prints the message “Hello, World”.

element interactivity effect (Sweller et al., 2011). The cognitive load imposed by element interactivity is said to be intrinsic to the material itself (**intrinsic cognitive load**), meaning that no matter how you present the material there will be a minimum level of difficulty.

How we present material can affect how difficult it is to understand it. For example, consider trying to teach beginner programmers what all of the keywords mean in a Java program that simply prints “Hello, World” to the screen (see Figure 1) and what each part of the code does. This would likely be too much to learn at once. Here there are many interacting elements (e.g. access modifiers, classes, types, functions, function arguments, comments, member fields, member functions, function calls, string literals) just in the code of this simple program. Compare this to simply pointing out the line that prints the message and the fact that anything between the quotation marks is printed to the terminal window. This approach is much easier to understand since there are fewer things that must be considered. Changing how the material is presented has the effect of increasing or reducing **extraneous cognitive load**, a type of cognitive load that is not intrinsic to the material itself but dependent on how the material is presented.

Following on with the example, the details of the “Hello, world” program could be taught more easily once the learner has learned the fundamentals and acquired a basic intuition for programming. It is theorised that this is because as complex ideas and tasks are learned, we learn more compact representations of these that are committed to long-term memory. These compact representations are often referred to as **schema**; a schema is essentially a single unit of knowledge representing a concept or task (Axelrod, 1973; Abelson, 1981; Bartlett et al., 1995). By acquiring a schema, a task with many interacting elements can be reduced to a task with a few, or even a single, element. And under the element interactivity effect, fewer interacting elements means less overall cognitive load.

To summarise CLT in a rather simplified manner, an individual has a limited amount of working memory to process information and if the amount of cognitive load exceeds the capacity of this working memory, then processing the information becomes much more difficult. There are three things from CLT that we are interested in: element interactivity, which is the main contributor to both intrinsic cognitive load and overall cognitive load; intrinsic cognitive load which is related to the inherent complexity of information; and extraneous load which is influenced by how the material/task is presented.

3 Conceptual Density in Text

To reiterate, conceptual density is an idea related to the degree to which concepts in a domain are integrated, and in this project, we are interested in conceptual density in the context of text documents. Some of the questions that we are interested in are, for example: how are concepts related, to what degree are concepts integrated, do concepts form self-referential systems, how are concepts referenced and presented in the document? Conceptual density can be rooted in element interactivity and described in terms of intrinsic cognitive load, and the way that concepts are presented in a document (e.g. how concepts are referenced between sections) can be described in terms of extraneous cognitive load. In this section, I will give a simple working definition of a concept in text and describe a few text features related to the idea of conceptual density.

3.1 Identifying Concepts in Text

To understand the relationships between concepts, we must first be able to identify concepts in text. A definition of a concept that could be used is a noun phrase - a phrase in text that denotes a thing. Using words from the previous sentence as an example, both “definition” and “simple definition” would be valid noun phrases; if we include determiners into the definition of a noun phrase “a simple definition” would also be a valid noun phrase. A more precise definition this type of noun phrase is: an optional determiner, followed by a (possibly empty) sequence of nouns and/or adjectives, terminated by a noun. This kind of noun phrase can also be expressed in regex-like syntax using part of speech tags:

Grammar 1 NP: <DT>?<NN.*|JJ>*<NN.*>

where the angle brackets denote a group/part of speech tag, DT is a determiner (e.g. a, the), NN.* is any type of noun and JJ is an adjective.

This pattern describes noun phrases such as ‘apple’, ‘an apple’ and ‘a red apple’. However, these noun phrases may also be joined together by prepositions (e.g. ‘in’ or ‘of’) or coordinating conjunctions (e.g. ‘and’ or ‘or’) to form a single, complex concept. For example, consider the sentence “Zeus is the sky and thunder god of the ancient Greek religion”. We can see that ‘the sky and thunder god’ is one concept because we are not saying that Zeus is ‘the sky’ and that he is separately the ‘thunder god’, so we have a case for considering two noun phrases joined by a coordinating conjunction as a single concept. It is also evident that we should consider ‘the sky and thunder god of the ancient Greek religion’ as a single, complex concept because Zeus is not the sky and thunder god of all religions, but rather of the ancient Greek religion alone. I call these types of noun phrases ‘complex noun phrases’. For this project, I define a concept to be a complex noun phrase, which is a noun phrase optionally followed by a sequence of preposition/coordinating conjunction and noun phrase pairs. We can define a complex noun phrase more precisely as:

Grammar 2 CNP: <NP>(<IN|CC><NP>)*

where NP is Grammar 1, IN is a preposition and CC is a coordinating conjunction.

3.2 Presentation of Concepts

The presentation of concepts is the way that concepts are talked about in a document and how concepts are referenced between sections of the document. For example, a concept in section A may require a reference to a concept in section B; likewise a concept in section C may require a reference to a concept in section B. These types of interactions between concepts are of interest, and are of relevance to the idea of conceptual density, since the interaction introduces another element that must be considered simultaneously by the reader. This increases element interactivity which is the core of cognitive load. In this section, I will propose two types of concepts present in text, and two types of references that could be used to help measure the effect of presentation on conceptual density.

3.2.1 A Priori and Emerging Concepts

It is possible to categorise concepts into two types: a priori and emerging. An a priori concept is a concept that the reader of a text document would be expected to know beforehand or to be common knowledge. An emerging concept is a concept that is defined or introduced within a document. For example, a textbook on computer programming would expect the reader to know what a computer is, but possibly not the concept of recursion. In this case, the concept ‘computer’ would be an a priori concept and ‘recursion’ would be an emerging concept. The importance of distinguishing between these two types of concepts is in how they contribute to cognitive load differently. Referencing an a priori concept would impose minimal cognitive load since the a priori concept would likely represent a fully acquired schema and well-integrated knowledge. On the other hand, an emerging concept would likely impose more cognitive load since it likely represents a schema that is yet to be fully acquired and still requires a noticeable amount of mental exertion in regards to the concept and its constituent parts.

3.2.2 Forward References

Forward references are where the text refers to a concept that is not fully explained until later in the document. These types of references introduce extraneous cognitive load since they make the reader *park* the involved concepts, without much existing knowledge to associate them with. Remembering unfamiliar terms that carry little meaning (to the reader) is more of a demanding task than remembering well-integrated knowledge. For example, when teaching Java programming to beginners the meaning and function of keywords such as `class` and `static` are often not explained and are usually just rote learned until later in the course.

3.2.3 Backward References

Backward references are where the text refers to a concept that was explained previously in the text. These types of references are likely to be introducing a relation between concepts. The cost of backward references is relatively low when compared to forward references. With forward references, we are making the readers consider an additional element that is not at all well understood by the reader, along with the current context. However, in the case of backward references, we are simply asking the reader to recall a previously

explained concept and possibly introduce a new relationship between the two concepts. This type of reference imposes low cognitive load as it requires the reader to call on pre-existing knowledge.

4 Building a Graph of Concepts

The domain of concepts present in a text document can be represented with a graph structure. The graph structure that is being built is similar to a mind-map, except instead of a single central idea/concept, we are building a rather free-form graph with concepts branching off other related concepts. Representing a domain of concepts in this type of graph has precedent in semantic networks/knowledge graphs. In this section, I will describe my implementation that currently works with text documents marked up in XML, how I build a graph structure, and how I extract concepts and the text features described in the previous section.

4.1 Nodes and Extracting Concepts

In the graph structure we are building, a node represents a concept present in a given text document. As discussed in Section 3.1, we create a definition of what constitutes a concept in terms of patterns of parts of speech. The formal definition of the pattern given defines what is called a grammar. The basic process to extract a concept from text is to segment the document into sentences, then split each sentence into a set of tokens (words separated by white-space), then for each set of tokens we assign a tag to each token denoting the part of speech of the token. From here we chunk the sentences into groupings based on the patterns described in the grammars. An example is shown in Figure 2.

By definition, a complex noun phrase is made up of multiple noun phrases (Grammar 2). For example, from the phrase “a dough of wheat flour and water” we could extract the noun phrases ‘a dough’, ‘wheat flour’ and ‘water’. And even from the noun phrase ‘wheat flour’ we can extract two distinct nouns, ‘wheat’ and ‘flour’. I choose to add these constituent parts to the graph as nodes to produce a more rich representation of the domain of concepts present in the text document. It also helps match up concepts that are closely related but perhaps not mentioned together or mentioned in the same form, e.g. dough’ and ‘bread dough’ in Figure 3a.

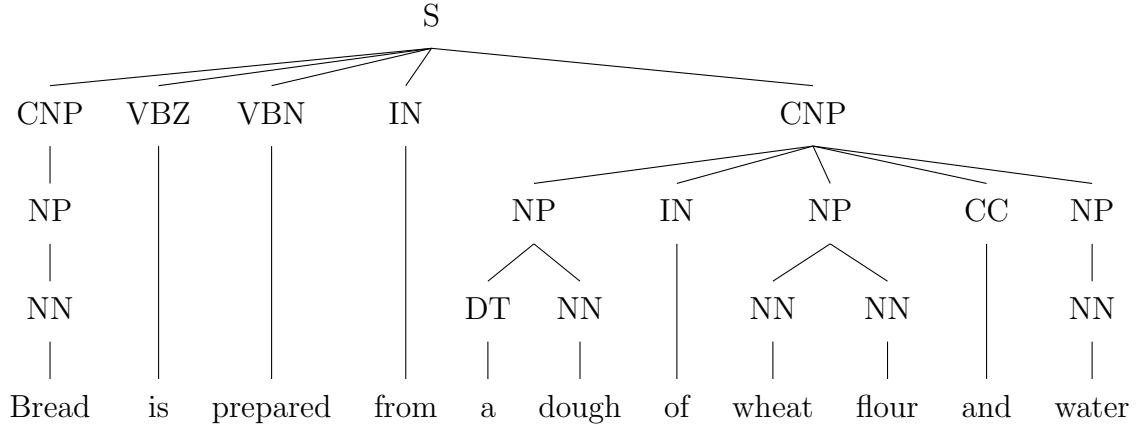


Figure 2: The parse tree resulting from first tagging the parts of speech for tokens in the sentence “Bread is prepared from a dough of wheat flour and water”, and then chunking the tagged tokens using the regex grammars defined in Grammar 1 and Grammar 2. The nodes directly above the terminal nodes denote each word’s part of speech tag. The tag ‘S’ represents a sentence, and the tags starting with ‘VB’ represent a verb.

4.2 Edges and Relating Concepts

In the graph structure we are building, a relation between concepts is represented as a directed edge between nodes. Edges are used to represent a type of relation where one concept refers to, or depends on, another concept. There are three criteria for which I choose to create an edge between two concepts: one concept is the subject of a sentence and the other concept appears in the same sentence; one concept is a complex noun phrase (Grammar 2) and the other is a noun phrase (Grammar 1) that is part of the complex noun phrase; or one concept is a noun phrase and the other is a noun or an adjective that is part of that noun phrase. The subject of a sentence can be identified through a dependency parse, which marks the relationships between words in a sentence (see Figure 4 for an example).

4.3 Identifying A Priori and Emerging Concepts

To identify a priori and emerging concepts we could look at how the concepts are mentioned. In particular, we could look for sentences that contain defi-

Bread

Bread

Bread is a staple food typically prepared from a dough of wheat flour and water, usually by baking.

Wheat

Wheat is a type of grain.

Wheat is commonly used for making wheat flour, a typical ingredient of bread dough.

(a) A simple text document on the topic of bread.

```
1 <document>
2   <title>
3     Bread.
4   </title>
5
6   <section>
7     <title>Bread</title>
8     <text>
9       Bread is a staple food typically prepared from a
10      dough of wheat flour and water, usually by baking.
11    </text>
12  </section>
13
14  <section>
15    <title>Wheat</title>
16    <text>
17      Wheat is a type of grain.
18      Wheat is commonly used for making wheat flour, a
19      typical ingredient of bread dough.
20    </text>
21  </section>
22 </document>
```

(b) A version of the document in Figure 3a that has been marked up in XML.

Figure 3: A sample document with a plain text version (Figure 3a) and a version marked up with XML (Figure 3b).




Figure 4: An example of a dependency parse.¹The arc labelled “nsubj” points to the subject of the sentence. For a description of what the relations on the arcs mean refer to (Martin and Jurafsky, 2018).

nitions. A type of sentence that could be considered as a type of definition could be one that follows the pattern ‘X is a Y’. If a concept appears as the subject of these kinds of sentences then perhaps that could indicate that the concept is an emerging concept. For example, if the first mention of ‘bread’ is “Bread is a staple food” then perhaps it is likely that bread is being defined here and thus is an emerging concept. A priori concepts could then be defined as any concept that does not fulfil the definition of an emerging concept.

It should be noted that this kind of approach would require knowledge of what kind of relations exist between concepts, a feature which has not been implemented yet in the current code. The current implementation works on the (incorrect) assumptions that: a concept that is only referenced from within a single section is an a priori concept; and concepts that are referenced from within multiple sections are emerging concepts. Implementing the method described in the previous paragraph is something to be looked at in future work.

4.4 Identifying Forward and Backward References

In the graph forward and backward references can be found through traversing the graph structure and identifying when the algorithm crosses between sections. A prerequisite is that we must record the sections that appear in the text, the order in which they are introduced during the parsing of the text document, and associate each node in the graph with a section. Each node is associated with the section that its corresponding concept appears in the most frequently. We can then identify a forward reference when the traversal

¹ Figure generated from <https://corenlp.run/>.

algorithm visits a node that is associated with a section that comes after the section of the previously visited node. Similarly, backward references can be identified when the traversal algorithm visits a node that is associated with a section that comes before the section of the previously visited node. Pseudocode for this process is given in Algorithm 1.

4.5 Deriving a Numerical Score

At the centre of conceptual density is the notion of the degree to which concepts are integrated. In a graph structure, this could be measured by looking at the connectivity of nodes. One useful measure may be the average outdegree:

$$S_D = \frac{1}{|V|} \sum_{v \in V} \deg^+(v) \quad (1)$$

where S_D is the score of the document D , V is the set of vertices in a graph, and $\deg^+(v)$ is the outdegree, or the number of outgoing edges from the vertex v . The interpretation of a high average degree is that the average concept is relatively difficult to learn since the concept is related to many other concepts which you must understand to understand the given concept.

We can further refine the scoring method by taking into account the strength of the edges, or in other words the weights of the edges:

$$S_D = \frac{1}{|V|} \sum_{v \in V} \sum_{e \in E[v]} e_w \quad (2)$$

where $E[v]$ is the set of edges originating from the vertex v , and e_w is the weight for a given edge e . The weight for a given edge can be adjusted based on the type of reference the edge represents. For example, we could set the weight of an edge to: 0.5 for edges that point to a priori concepts since they typically represent a concept that the reader is expected to know beforehand and are easier to process; 1.5 for backward references to emerging concepts since the reader is required to recall a concept that has been introduced within the text and using a poorly integrated schema will require more mental exertion than when considering an a priori concept; 2.0 for forward references since the reader is required to *park* a new concept without anything concrete to pin it too, increasing cognitive load; and 1.0 for any other edges that do not fall into any of these categories. The exact numbers are not important, but rather the relative scale of these numbers. References to a priori concepts

Algorithm 1: Marking Forward and Backward References

Data: *sections*: List of sections in a given document D

nodes: Set of nodes in the graph G

adj: Adjacency list for a directed graph G .

```
1
2 visited =  $\emptyset$ 
3
4 for node  $\in$  nodes do
5   | markEdges(node, NULL, visited)
6 end
7
8 Function markEdges(curr, prev, visited)
   | /* Arguments                                     */
   | /* curr: Current node                             */
   | /* prev: The previously visited node               */
   | /* visited: Set of visited nodes                   */
9
10  if prev  $\neq$  NULL AND curr.section  $\neq$  prev.section then
   | /* Find the ordering of the sections associated with
   |    the nodes curr and prev.                                     */
11   | curr_i  $\leftarrow$  sections.indexOf(curr.section)
12   | prev_i  $\leftarrow$  sections.indexOf(prev.section)
13
14   | if curr_i < prev_i then
15   |   | Mark edge {prev, curr} as a backward reference
16   | else if curr_i > prev_i then
17   |   | Mark edge {prev, curr} as a forward reference
18
19   | if curr  $\notin$  visited then
20   |   | visited  $\leftarrow$  visited  $\cup$  {curr}
21
22   | for neighbour  $\in$  adj[curr] do
23   |   | markEdges(neighbour, curr, visited)
24   | end
25 end
```

should impose less cognitive load than backward references, and backward references should impose less cognitive load than forward references.

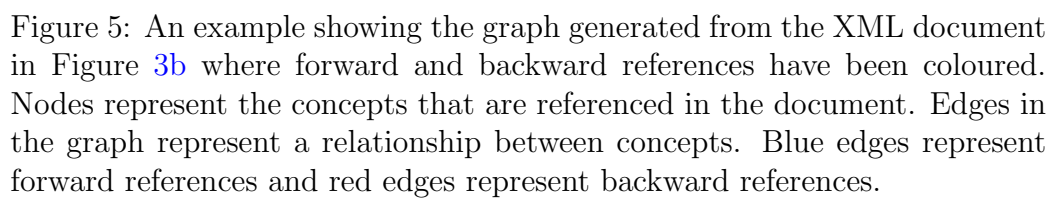
4.6 Summary

By using part of speech tagging and a regex-based chunker we can extract concepts from text. We can then derive relationships between concepts by looking at dependency parses of sentence structures and by looking at the constituent parts of complex concepts represented by noun phrases. For the former method, we simply take the subject of a sentence and create edges between it and each of the other concepts that appear in the same sentence. Concepts are represented as a node in the graph and relationships between concepts are represented as edges in the graph. The different types of concepts and references can be identified by analysing the edges in the graph. Finally, we can derive a score from the generated graph by measuring the average weighted outdegree. Figure 5 shows the graph generated using the current implementation and the sample text document in Figure 3.

5 Related Work

5.1 Information Extraction

The process of identifying concepts and how they are related that I have described in this report is similar to the task of information extraction. However, there are some key differences between what I have done and what is usually done in information extraction. First and foremost, information extraction seems to typically be concerned with named-entity recognition - the identification of named-entities such as names of people, companies, locations, and dates. In this project, we are concerned with concepts in general (all types of entities), which can be thought of as a superset of named-entities. Furthermore, information extraction is also interested in the relationship between named-entities. However, in this project and the work-to-date, we are not too concerned about the nature of a relationship between concepts as much as the fact that a relationship exists.



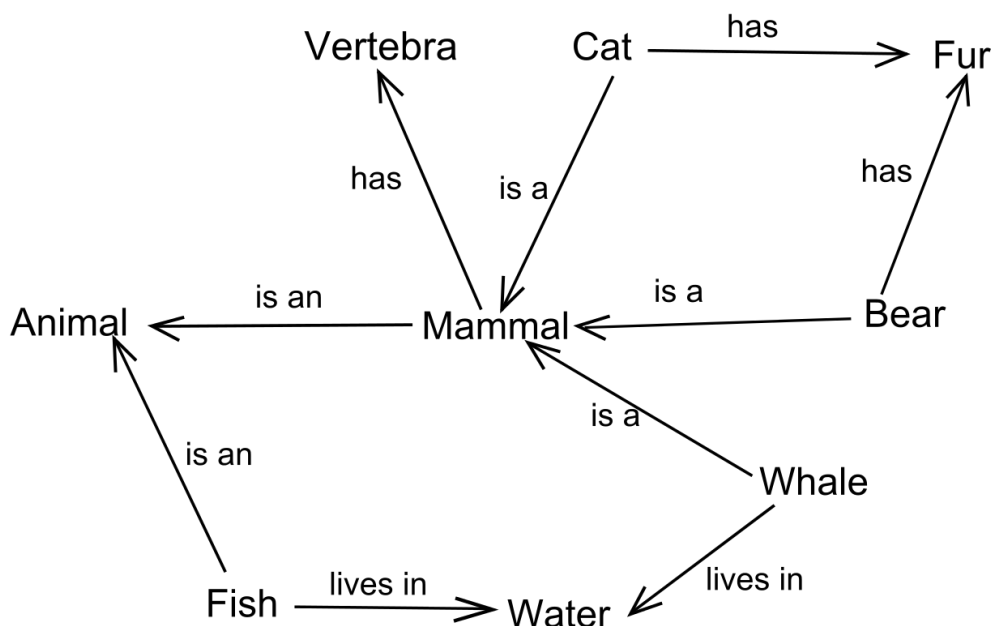


Figure 6: An example of a semantic network.²

5.2 Semantic Networks and Knowledge Graphs

The type of graph structure that is being built up in this project is very similar to semantic networks/knowledge graphs. Some differences between the graphs in this project and semantic networks are that: we are not annotating edges with the type of relationship (compare Figure 5 and Figure 6, for example), and the criteria for deciding what constitutes a node or edge in the graph may differ from what is common in the literature. For one, somewhat complex noun phrases are added as nodes in this project rather than just simple noun phrases or the head of noun phrases.

²Figure from https://commons.wikimedia.org/wiki/File:Semantic_Net.svg.

6 Future Work

6.1 Transitioning to Unstructured Text

In the current implementation a graph can be built from semi-structured text (e.g. Figure 5), but recreating the same graph from the corresponding plain text version of a document remains difficult. Moving to unstructured text is desirable since using semi-structured text requires someone to manually mark up the text, which is a time-consuming process. At present, the sections and section headers must be explicitly marked up in XML. Identifying the section titles and where sections start and end automatically can be unreliable at times when using simple regex pattern matching due to different section heading formats and it is liable to false positives (see Table 1 in the appendix, for example).

6.2 Improving Information Extraction

One issue with the current implementation is the quality of the extracted concepts and how the concepts are related. There are many cases where the extracted concepts and relations do not align with what a human would intuitively pick out. For example, under the current set of grammars, concepts that are captured by noun phrases are well covered but many concepts are not nouns but ‘doing things’, in other words, verbs or verb phrases. For example, in programming, we can talk about ‘calling functions’, ‘looping’ and ‘instantiating objects’ as concepts, and all of these are examples of verbs or verb phrases. At this point, the only type of verb that I have included in the current implementation is the gerund verb, a type of verb that is typically suffixed with -ing, such as **eating**, **sleeping**. While the current implementation is adequate to build up a graph structure, important concepts and relations between concepts are likely to be missed. So it is important that future work focuses on improving the quality of extracted concepts and relations by: extending the definition of a concept in text to include verbs and verb phrases; and improving the quality of the extracted relations between concepts.

There is an OpenIE (Open Information Extraction) framework that has been implemented by the NLP group at Stanford University ([Manning et al., 2014](#); [Angeli et al., 2015](#)) which is made for performing information extraction, which is essentially the task that I am trying to perform. The benefits

of using this is that many of the issues with the current implementation for extracting concepts and relations between concepts in this project are likely to have been solved. One issue with using Stanford’s OpenIE Implementation is that it generates many spurious relation triples (see Table 2 in the appendix for an illustrative example). So if this framework were to be used in this project the main thing that would need to be worked on is filtering the generated relation triples.

6.3 Additional Graph Features

Many graph analysis techniques have not yet been explored in this project. For example, see (NetworkX Developers, 2004). It would be interesting to see if incorporating features derived from these analysis techniques into the scoring algorithm would improve it. For example, we could look at: identifying simple cycles in the graph using Johnson’s algorithm; identifying minimum cuts which could be used to identify bottleneck concepts and to find sub-graphs that are *isolated* in the sense that they are mostly separate from the main supergraph; and reachability of nodes in the graph.

6.4 Additional Text Features

There are other features that could be extracted from text. In (Robins, 2010) it is suggested that a measure of proximity of key concept terms could be used to help measure conceptual density. Some methods that may help with this are: various keyword extraction techniques (Mihalcea and Tarau, 2004; Matsuo and Ishizuka, 2004; Ercan and Cicekli, 2007; Rose et al., 2010); and long-distance dependency as a metric of reading comprehension difficulty (Liu, 2008).

6.5 Evaluation of the Model

The quality of the model for quantifying conceptual density has not yet been evaluated. Since no objective measure of conceptual density exists in the current literature, there exists no benchmarks to compare the model with. The next best thing to compare the model against is human judgements. One experiment that could be done is to gather a collection of snippets from textbooks. These snippets would each consist of a few chapters from a given textbook. Then we could arrange the snippets into pairs and get

human judgements on whether or not the first snippet is more conceptually dense than the second snippet. From there, for each document pair, we could generate scores using the model described in this report and determine which document the model judges as being more conceptually dense. Then we could compare the model’s judgements to the human judgements to evaluate how ‘good’ the model is, or how well it follows human intuition.

7 Conclusion

In this report, I have explored a graph-based approach to quantifying conceptual density - an idea relating to the degree that concepts of a particular knowledge domain, encoded by a text document, are integrated. I have discussed how this idea of conceptual density is grounded in, and related to, learning edge momentum and cognitive load theory. I have discussed how the notion of conceptual density can be defined and measured in text documents. By analysing the language in the documents and observing the patterns in natural language we can create an operational definition of a ‘concept’ in text. Then by analysing the relationships between concepts in a given document we can build up a graph structure from which we can derive a score of conceptual density.

References

- Robert P Abelson. Psychological status of the script concept. *American psychologist*, 36(7):715, 1981.
- Gabor Angeli, Melvin Jose Johnson Premkumar, and Christopher D Manning. Leveraging linguistic structure for open domain information extraction. In *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 344–354, 2015.
- Robert Axelrod. Schema theory: An information processing model of perception and cognition. *American political science review*, 67(4):1248–1266, 1973.

- Frederic Charles Bartlett, Frederic C Bartlett, and Walter Kintsch. Remembering: A study in experimental and social psychology, 1995.
- Paul Chandler and John Sweller. Cognitive load while learning to use a computer program. *Applied cognitive psychology*, 10(2):151–170, 1996.
- Gonenc Ercan and Ilyas Cicekli. Using lexical chains for keyword extraction. *Information Processing & Management*, 43(6):1705–1714, 2007.
- Rik Koncel-Kedziorski, Dhanush Bekal, Yi Luan, Mirella Lapata, and Hananeh Hajishirzi. Text generation from knowledge graphs with graph transformers. *arXiv preprint arXiv:1904.02342*, 2019.
- Haitao Liu. Dependency distance as a metric of language comprehension difficulty. *Journal of Cognitive Science*, 9(2):159–191, 2008.
- Christopher D. Manning, Mihai Surdeanu, John Bauer, Jenny Finkel, Steven J. Bethard, and David McClosky. The Stanford CoreNLP natural language processing toolkit. In *Association for Computational Linguistics (ACL) System Demonstrations*, pages 55–60, 2014. URL <http://www.aclweb.org/anthology/P/P14/P14-5010>.
- James H Martin and Daniel Jurafsky. *Speech and language processing: An introduction to natural language processing, computational linguistics, and speech recognition. Third Edition draft*. Pearson/Prentice Hall Upper Saddle River, 2018.
- Yutaka Matsuo and Mitsuru Ishizuka. Keyword extraction from a single document using word co-occurrence statistical information. *International Journal on Artificial Intelligence Tools*, 13(01):157–169, 2004.
- Rada Mihalcea and Paul Tarau. Textrank: Bringing order into text. In *Proceedings of the 2004 conference on empirical methods in natural language processing*, pages 404–411, 2004.
- NetworkX Developers. Algorithms - networkx 2.3 documentation. <https://networkx.github.io/documentation/stable/reference/algorithms/index.html>, 2004. Accessed on 05 July 2019.
- Anthony Robins. Learning edge momentum: A new account of outcomes in cs1. *Computer Science Education*, 20(1):37–71, 2010.

- Stuart Rose, Dave Engel, Nick Cramer, and Wendy Cowley. Automatic keyword extraction from individual documents. *Text mining: applications and theory*, 1:1–20, 2010.
- John F Sowa. Semantic networks. 1987.
- John Sweller. Cognitive load theory, learning difficulty, and instructional design. *Learning and instruction*, 4(4):295–312, 1994.
- John Sweller, Paul Ayres, and Slava Kalyuga. The element interactivity effect. In *Cognitive Load Theory*, pages 193–201. Springer, 2011.
- Liecai Zhang. *Knowledge graph theory and structural parsing*. Twente University Press, 2002.

A Tables

Table 1: An example of the current regex-based implementation for finding section headers and where sections start and end, using the text of the paper by [Sweller et al. \(2011\)](#). Asterisks mark cases where portions of text that are not section headers have been erroneously picked out.

Extracted Section Header
Chapter 15
The Element Interactivity Effect
Empirical Evidence for the Element Interactivity Effect
*The secondary task was presented on a separate computer and involved a tone
*When instructional materials involved low or no interaction between elements of
Element Interactivity and Understanding Instructions
*The extent to which we understand instructions depends on levels of element
Element Interactivity and the Modality Effect
Element Interactivity and the Expertise Reversal Effect
Element Interactivity and the Imagination Effect
Conditions of Applicability

Extracted Section Header
*There is also an obvious relation between these two effects as the levels of element interactivity that produce an intrinsic cognitive load are always relative to levels of
Instructional Implications
Conclusion

Table 2: An example of the relation triples extracted from the document in Figure 3a illustrating how the OpenIE implementation in Stanford CoreNLP generates many spurious triples. Notice that some of the relation triples are nonsensical, e.g. (‘Bread’, ‘is’, ‘staple food prepared’) in the second row.

Entity 1	Relation	Entity 2
Bread	is	food typically prepared from dough by baking
Bread	is	staple food prepared
Bread	is	staple food typically prepared from dough of wheat flour usually by baking
Bread	is	food typically prepared from dough
Bread	is	food by baking
Bread	is	staple food typically prepared from dough usually by baking
Bread	is	staple food prepared from dough by baking
Bread	is	food typically prepared usually by baking
Bread	is	food prepared
Bread	is	food typically prepared
Bread	is	staple
Bread	is	food prepared from dough of wheat flour by baking
Bread	is	staple food typically prepared
Bread	is	staple food by baking
Bread	is	food prepared from dough by baking
Bread	is	staple food typically prepared by baking
Bread	is	staple food prepared from dough of wheat flour by baking
Bread	is	food typically prepared from dough of wheat flour

Entity 1	Relation	Entity 2
Bread	is	staple food
Bread	is	food typically prepared by baking
Bread	is	staple food typically prepared from dough of wheat flour by baking
Bread	is	food typically prepared from dough of wheat flour by baking
Bread	is	staple food typically prepared from dough
Bread	is	food prepared from dough
Bread	is	food prepared from dough of wheat flour
Bread	is	staple food typically prepared usually by baking
Bread	is	staple food prepared from dough usually by baking
Bread	is	staple food prepared from dough of wheat flour usually by baking
Bread	is	food by usually baking
Bread	is	food prepared usually by baking
Bread	is	food prepared by baking
Bread	is	staple food by usually baking
Bread	is	food typically prepared from dough of wheat flour usually by baking
Bread	is	food prepared from dough usually by baking
Bread	is	staple food prepared usually by baking
Bread	is	staple food prepared by baking
Bread	is	staple food typically prepared from dough by baking
Bread	is	staple food typically prepared from dough of wheat flour
Bread	is	staple food prepared from dough of wheat flour
Bread	is	food
Bread	is	food typically prepared from dough usually by baking
Bread	is	food prepared from dough of wheat flour usually by baking
Bread	is	staple food prepared from dough

Entity 1	Relation	Entity 2
Wheat	making	ingredient of bread dough
Wheat	making	wheat flour
Wheat	making	ingredient
Wheat	is	commonly used
wheat flour	ingredient of	bread dough
Wheat	making	typical ingredient of bread dough
Wheat	making	typical ingredient
Wheat	is	used

B Aims and Objectives

The project is still on track in regards to the original aims and objectives and they have not been revised.

Aims Conceptual Density describes the degree to which concepts in an extended text (such as the textbook for a university course) are interdependent, or to what degree they interact with each other. There is a hypothesis that text documents with high conceptual density are harder for students to process, because of the high interdependence between concepts. We aim to establish a method of quantifying this idea of conceptual density in text corpora.

Objectives

- Background reading on:
 - Conceptual Density
 - Cognitive Load
 - Hypertext Markers
 - Textual Signposting
 - Topic Segmentation
 - Other relevant topics
- Identify suitable text corpora
- Design and implement experiments to test published natural language processing methods on the above corpora
- If time allows, develop, implement, and test new methods for quantifying conceptual density.

Timeline

- Mar-Apr: Background reading
- Apr-May: Identifying text corpora
- Jun: Designing experiments & preliminary results (Interim report)
- Jul-Aug: Implementing experiments
- Sep-Oct: Developing, implementing, and testing new methods for quantifying conceptual density (Final report).